

CS3315 Project: MLP with Keras

```
In [1]: import numpy as np
import pandas as pd
from sklearn.linear_model import Perceptron
```

```
In [2]: # import data
filename = 'data/2020.06.19.csv'
df = pd.read_csv(filename)

# sample small subset
# df = df.sample(500000, random_state=78)
df.info()
df.head(2)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 765360 entries, 0 to 765359
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   avg_ipt                765360 non-null float64
1   bytes_in               765360 non-null int64  
2   bytes_out              765360 non-null int64  
3   dest_ip                765360 non-null int64  
4   dest_port              740863 non-null float64
5   entropy                765360 non-null float64
6   num_pkts_out           765360 non-null int64  
7   num_pkts_in            765360 non-null int64  
8   proto                  765360 non-null int64  
9   src_ip                 765360 non-null int64  
10  src_port               740863 non-null float64
11  time_end               765360 non-null int64  
12  time_start             765360 non-null int64  
13  total_entropy          765360 non-null float64
14  label                  765360 non-null object
15  duration               765360 non-null float64
dtypes: float64(6), int64(9), object(1)
memory usage: 93.4+ MB
```

Out[2]:

	avg_ipt	bytes_in	bytes_out	dest_ip	dest_port	entropy	num_pkts_out	num_pkts_in	protc
0	7.5	342	3679	786	9200.0	5.436687	2	2	€
1	0.0	0	0	786	55972.0	0.000000	1	1	€

```
In [3]: # clean data
df.dropna(inplace=True)
df.isna().sum()

# need to clean for features that are 0 and don't make sense (bytes = 0)
```

```
Out[3]: avg_ipt      0
        bytes_in    0
        bytes_out    0
        dest_ip      0
        dest_port    0
        entropy      0
        num_pkts_out  0
        num_pkts_in  0
        proto        0
        src_ip       0
        src_port     0
        time_end     0
        time_start   0
        total_entropy 0
        label        0
        duration     0
        dtype: int64
```

```
In [4]: print('label values:', df['label'].unique())

def ordinal_encoder(category):
    dict = {'benign':0, 'outlier':1, 'malicious':2}
    return dict[category]

print('benign', ordinal_encoder('benign'))
print('outlier', ordinal_encoder('outlier'))
print('malicious', ordinal_encoder('malicious'))
df['label'] = df['label'].apply(ordinal_encoder)

label values: ['benign' 'outlier' 'malicious']
benign 0
outlier 1
malicious 2
```

```
In [5]: features = ['avg_ip',
                    'bytes_in',
                    'bytes_out',
                    'dest_ip',
                    'dest_port',
                    'entropy',
                    'num_pkts_in',
                    'num_pkts_out',
                    'proto',
                    'src_ip',
                    'src_port',
                    'time_end',
                    'time_start',
                    'total_entropy',
                    'duration']

X = df.loc[:, features]
y = df.loc[:, 'label']
```

```
In [6]: # Scale features
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
# try PolyScaler?

scaler = StandardScaler()
scaler.fit(X)
X = scaler.transform(X)
```

```
In [7]: # test/train split

from sklearn.model_selection import train_test_split

# 80/20 training/validation split
X_train, X_val, y_train, y_val = train_test_split(X,y, train_size=.8, t

y_train = y_train.to_numpy()
y_val = y_val.to_numpy()

# should print number of shape: (num features, num entries)
print('Training set: ', 'X: ', X_train.shape, 'y: ', y_train.shape, 'Va

print(X_train[1])
print(y_train[1])
```

```
Training set: X: (592690, 15) y: (592690,) Validation set: X: (14
8173, 15) printy: (148173,)
[-0.05565213 -0.24851657 -0.4012423 -0.13191521 -0.1413975 -1.344987
3
-0.27015576 -0.2850236 -0.40847424 -0.50115682 0.54239382 0.332459
24
0.33237538 -0.20292837 -0.33436206]
0
```

In [8]: *# import tensorflow and keras*

```
import tensorflow as tf
from tensorflow import keras
print(tf.__version__)
print(keras.__version__)

import os
```

2022-12-01 19:22:25.620736: I tensorflow/core/platform/cpu_feature_guard.cc:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: AVX2 FMA
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

2.12.0-dev20221112
2.12.0

In [9]: *# keras sequential API model (pg 299)*

```
model = keras.models.Sequential([
    keras.layers.InputLayer(input_shape=(X_train.shape[1])),
    keras.layers.Dense(300, activation="relu"),
    keras.layers.Dense(200, activation="relu"),
    keras.layers.Dense(3, activation="softmax")
])
print(X_train.shape[1])

model.summary()
```

15
Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 300)	4800
dense_1 (Dense)	(None, 200)	60200
dense_2 (Dense)	(None, 3)	603

=====
Total params: 65,603
Trainable params: 65,603
Non-trainable params: 0
=====

In [10]: *# TODO: adjust kernel initializer or bias initializer?*
<https://keras.io/initializers/>

```
In [11]: # compile the model
model.compile(loss="sparse_categorical_crossentropy",
              optimizer = "sgd",
              metrics="accuracy")
```

```

In [12]: # train the model
num_epochs=30
print(X_train.shape)
history = model.fit(X_train, y_train,
                    epochs=num_epochs,
                    validation_data=(X_val, y_val))

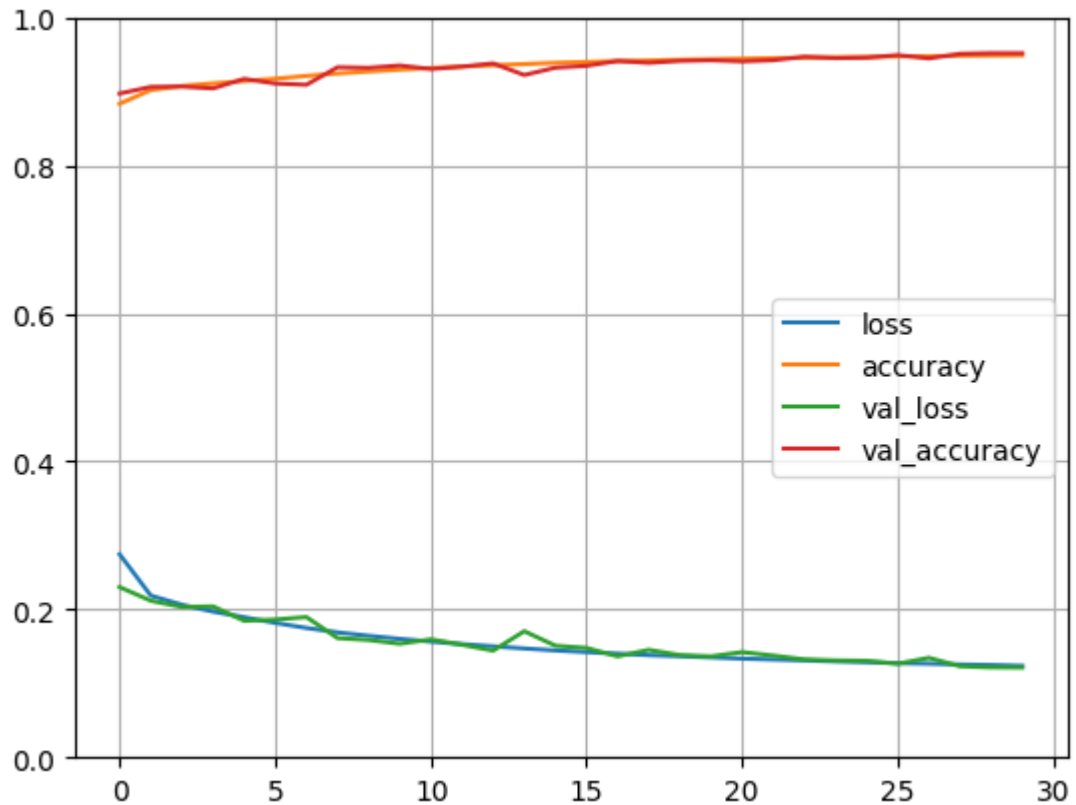
# loss is nan -> probably indicative of exploding gradients -- try again

(592690, 15)
Epoch 1/30
18522/18522 [=====] - 35s 2ms/step - loss: 0.
2745 - accuracy: 0.8832 - val_loss: 0.2302 - val_accuracy: 0.8971
Epoch 2/30
18522/18522 [=====] - 36s 2ms/step - loss: 0.
2188 - accuracy: 0.9018 - val_loss: 0.2118 - val_accuracy: 0.9061
Epoch 3/30
18522/18522 [=====] - 33s 2ms/step - loss: 0.
2063 - accuracy: 0.9073 - val_loss: 0.2033 - val_accuracy: 0.9065
Epoch 4/30
18522/18522 [=====] - 32s 2ms/step - loss: 0.
1971 - accuracy: 0.9108 - val_loss: 0.2041 - val_accuracy: 0.9040
Epoch 5/30
18522/18522 [=====] - 33s 2ms/step - loss: 0.
1895 - accuracy: 0.9140 - val_loss: 0.1847 - val_accuracy: 0.9169
Epoch 6/30
18522/18522 [=====] - 33s 2ms/step - loss: 0.
1818 - accuracy: 0.9171 - val_loss: 0.1864 - val_accuracy: 0.9107
Epoch 7/30
18522/18522 [=====] - 33s 2ms/step - loss: 0.
1750 - accuracy: 0.9209 - val_loss: 0.1897 - val_accuracy: 0.9089
Epoch 8/30
18522/18522 [=====] - 32s 2ms/step - loss: 0.
1690 - accuracy: 0.9243 - val_loss: 0.1613 - val_accuracy: 0.9326
Epoch 9/30
18522/18522 [=====] - 33s 2ms/step - loss: 0.
1644 - accuracy: 0.9270 - val_loss: 0.1587 - val_accuracy: 0.9319
Epoch 10/30
18522/18522 [=====] - 33s 2ms/step - loss: 0.
1602 - accuracy: 0.9294 - val_loss: 0.1537 - val_accuracy: 0.9350
Epoch 11/30
18522/18522 [=====] - 32s 2ms/step - loss: 0.
1566 - accuracy: 0.9316 - val_loss: 0.1596 - val_accuracy: 0.9303
Epoch 12/30
18522/18522 [=====] - 33s 2ms/step - loss: 0.
1531 - accuracy: 0.9336 - val_loss: 0.1523 - val_accuracy: 0.9334
Epoch 13/30
18522/18522 [=====] - 32s 2ms/step - loss: 0.
1502 - accuracy: 0.9354 - val_loss: 0.1443 - val_accuracy: 0.9378
Epoch 14/30
18522/18522 [=====] - 34s 2ms/step - loss: 0.
1473 - accuracy: 0.9369 - val_loss: 0.1706 - val_accuracy: 0.9221
Epoch 15/30
18522/18522 [=====] - 34s 2ms/step - loss: 0.
1446 - accuracy: 0.9386 - val_loss: 0.1510 - val_accuracy: 0.9321
Epoch 16/30

```

18522/18522 [=====] - 32s 2ms/step - loss: 0.
1424 - accuracy: 0.9398 - val_loss: 0.1479 - val_accuracy: 0.9346
Epoch 17/30
18522/18522 [=====] - 32s 2ms/step - loss: 0.
1403 - accuracy: 0.9410 - val_loss: 0.1366 - val_accuracy: 0.9411
Epoch 18/30
18522/18522 [=====] - 32s 2ms/step - loss: 0.
1382 - accuracy: 0.9419 - val_loss: 0.1454 - val_accuracy: 0.9386
Epoch 19/30
18522/18522 [=====] - 32s 2ms/step - loss: 0.
1368 - accuracy: 0.9428 - val_loss: 0.1382 - val_accuracy: 0.9412
Epoch 20/30
18522/18522 [=====] - 32s 2ms/step - loss: 0.
1352 - accuracy: 0.9435 - val_loss: 0.1366 - val_accuracy: 0.9421
Epoch 21/30
18522/18522 [=====] - 32s 2ms/step - loss: 0.
1335 - accuracy: 0.9442 - val_loss: 0.1423 - val_accuracy: 0.9406
Epoch 22/30
18522/18522 [=====] - 32s 2ms/step - loss: 0.
1323 - accuracy: 0.9450 - val_loss: 0.1376 - val_accuracy: 0.9421
Epoch 23/30
18522/18522 [=====] - 32s 2ms/step - loss: 0.
1312 - accuracy: 0.9455 - val_loss: 0.1324 - val_accuracy: 0.9468
Epoch 24/30
18522/18522 [=====] - 32s 2ms/step - loss: 0.
1298 - accuracy: 0.9461 - val_loss: 0.1311 - val_accuracy: 0.9450
Epoch 25/30
18522/18522 [=====] - 32s 2ms/step - loss: 0.
1287 - accuracy: 0.9466 - val_loss: 0.1306 - val_accuracy: 0.9455
Epoch 26/30
18522/18522 [=====] - 32s 2ms/step - loss: 0.
1274 - accuracy: 0.9474 - val_loss: 0.1266 - val_accuracy: 0.9490
Epoch 27/30
18522/18522 [=====] - 32s 2ms/step - loss: 0.
1268 - accuracy: 0.9476 - val_loss: 0.1346 - val_accuracy: 0.9446
Epoch 28/30
18522/18522 [=====] - 32s 2ms/step - loss: 0.
1256 - accuracy: 0.9482 - val_loss: 0.1232 - val_accuracy: 0.9508
Epoch 29/30
18522/18522 [=====] - 32s 2ms/step - loss: 0.
1249 - accuracy: 0.9485 - val_loss: 0.1217 - val_accuracy: 0.9516
Epoch 30/30
18522/18522 [=====] - 32s 2ms/step - loss: 0.
1238 - accuracy: 0.9489 - val_loss: 0.1215 - val_accuracy: 0.9516

```
In [13]: # plot loss vs. accuracy (HOML p. 305)
import matplotlib.pyplot as plt
pd.DataFrame(history.history).plot()
plt.grid(True)
plt.gca().set_ylim(0,1)
plt.show()
```



```
In [14]: # save model
model.save("project_mlp_2_layers")
```

INFO:tensorflow:Assets written to: project_mlp_2_layers/assets

```
In [15]: # test predictions
X_new = X_val
test_pred = np.argmax(model.predict(X_new), axis=-1)
```

4631/4631 [=====] - 4s 754us/step


```
In [16]: from sklearn.metrics import *
print("Predicted labels:\t", test_pred)
print("Actual labels:\t\t", y_val)
print(classification_report(y_val, test_pred))
```

```
Predicted labels:      [0 1 0 ... 2 2 2]
Actual labels:         [0 1 0 ... 2 2 1]

              precision    recall  f1-score   support

     0           1.00        1.00        1.00       75511
     1           0.91        0.80        0.85       24572
     2           0.90        0.96        0.93       48090

 accuracy              0.95       148173
 macro avg           0.94        0.92        0.93       148173
 weighted avg        0.95        0.95        0.95       148173
```

Validate Model with Data from June 2022

```
In [17]: # import data
filename = 'data/2020.06.20.csv'
df2 = pd.read_csv(filename)

# sample small subset
#df2 = df2.sample(n=100000, random_state=78)
df2.info()
df2.head(2)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 770853 entries, 0 to 770852
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   avg_ipt                770853 non-null float64
1   bytes_in               770853 non-null int64
2   bytes_out              770853 non-null int64
3   dest_ip                770853 non-null int64
4   dest_port              770853 non-null int64
5   entropy                770853 non-null float64
6   num_pkts_out           770853 non-null int64
7   num_pkts_in            770853 non-null int64
8   proto                  770853 non-null int64
9   src_ip                 770853 non-null int64
10  src_port                770853 non-null int64
11  time_end                770853 non-null int64
12  time_start              770853 non-null int64
13  total_entropy           770853 non-null float64
14  label                   770853 non-null object
15  duration                770853 non-null float64
dtypes: float64(4), int64(11), object(1)
memory usage: 94.1+ MB
```

Out[17]:

	avg_ipt	bytes_in	bytes_out	dest_ip	dest_port	entropy	num_pkts_out	num_pkts_in	pro
0	34.57143	34	29	786	5900	5.040459	7	10	
1	37.00000	34	29	786	5900	5.127916	7	10	

```
In [18]: # clean data
df2.dropna(inplace=True)
df2.isna().sum()
```

```
Out[18]: avg_ip      0
bytes_in    0
bytes_out   0
dest_ip     0
dest_port   0
entropy     0
num_pkts_out 0
num_pkts_in 0
proto       0
src_ip      0
src_port    0
time_end    0
time_start  0
total_entropy 0
label       0
duration    0
dtype: int64
```

```
In [19]: print('label values:', df2['label'].unique())

def ordinal_encoder(category):
    dict = {'benign':0, 'outlier':1, 'malicious':2}
    return dict[category]

print('benign', ordinal_encoder('benign'))
print('outlier', ordinal_encoder('outlier'))
print('malicious', ordinal_encoder('malicious'))
df2['label'] = df2['label'].apply(ordinal_encoder)

label values: ['malicious' 'benign' 'outlier']
benign 0
outlier 1
malicious 2
```

```
In [20]: features = ['avg_ip',
                    'bytes_in',
                    'bytes_out',
                    'dest_ip',
                    'dest_port',
                    'entropy',
                    'num_pkts_in',
                    'num_pkts_out',
                    'proto',
                    'src_ip',
                    'src_port',
                    'time_end',
                    'time_start',
                    'total_entropy',
                    'duration']

X_22 = df2.loc[:, features]
y_22 = df2.loc[:, 'label']
```

```
In [21]: # Scale features
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
# try PolyScaler?

scaler = StandardScaler()
scaler.fit(X_22)
X_22 = scaler.transform(X_22)

# change labels to numpy
y_22 = y_22.to_numpy()
```

```
In [22]: # test predictions
X_test_new = X_22
test_pred_22 = np.argmax(model.predict(X_test_new), axis=-1)

24090/24090 [=====] - 18s 767us/step
```

```
In [23]: print("Predicted labels:\t", test_pred_22)
print("Actual labels:\t\t", y_22)
print(classification_report(y_22, test_pred_22))
```

Predicted labels:	[2 2 2 ... 2 2 2]			
Actual labels:	[2 2 2 ... 1 1 2]			
	precision	recall	f1-score	support
0	1.00	0.99	1.00	366310
1	0.28	0.26	0.27	69389
2	0.84	0.86	0.85	335154
accuracy			0.87	770853
macro avg	0.71	0.70	0.71	770853
weighted avg	0.87	0.87	0.87	770853

In []: